

# Bitwise Operations

Instructor: Jeeho Ryoo

# Bitwise operations

Logical NOT: !

A	!A
T	F
F	T

# Bitwise operations

Logical NOT: !

A	!A
non-zero	0
zero	1

# Bitwise operations

Bitwise NOT:  $\sim$

A	$\sim A$
00	11
01	10
10	01
11	00

# Bitwise operations

Logical AND: &&

A	B	A && B
T	T	T
T	F	F
F	T	F
F	F	F

# Bitwise operations

Logical AND: &&

A	B	A && B
non-zero	non-zero	1
non-zero	zero	0
zero	non-zero	0
zero	zero	0

# Bitwise operations

Bitwise AND: &

A	B	A & B
1	1	1
1	0	0
0	1	0
0	0	0

Example:  $1100 \& 0101 = 0100$

# Bitwise operations

Logical OR: `||`

A	B	A    B
T	T	T
T	F	T
F	T	T
F	F	F



# Bitwise operations

Logical OR: `||`

A	B	A    B
non-zero	non-zero	1
non-zero	zero	1
zero	non-zero	1
zero	zero	0

# Bitwise operations

Bitwise OR: |

A	B	A   B
1	1	1
1	0	1
0	1	1
0	0	0

Example:  $1100 \mid 0101 = 1101$

# Bitwise operations

Logical XOR:

A	B	A XOR B
T	T	F
T	F	T
F	T	T
F	F	F

# Bitwise operations

Bitwise XOR: ^

A	B	A ^ B
1	1	0
1	0	1
0	1	1
0	0	0

Example:  $1100 \wedge 0101 = 1001$

## Bitwise operations

### Exercise:

Can you write a function that checks whether the third bit (from right) of a number is 1 or 0? Signature should be:

```
int thirdBitFromRight(int n);
```

# Bitwise operations

## Example:

```
#include <stdio.h>
int thirdBitFromRight(int n) {
    int mask = 4;
    return (n & mask) == 4;
}
void runTest(int n) {
    printf("n = %d, thirdBitFromRight = %d\n", n,
thirdBitFromRight(n));
}
int main() {
    runTest(4);
    runTest(15);
    runTest(0);
    runTest(11);
    runTest(-1);
    return 0;
}
```

# Bitwise operations

## Bit Masking:

A bit mask is an integer whose binary representation is intended to combine with another value using `&`, `|` or `^` to extract or set a particular bit or set of bits.

For example mask = 4 in the code from previous slide.

## Bitwise operations

Another exercise:

Write a function that turns “on” the third bit (from right):



## Practice Problems

**MEDIUM** – Write a function that takes a number and turns **on** the first and third binary digits (from right) for this number. Here are some examples:

$$8 = (1000)_2 \rightarrow 13 = (1101)_2$$

$$0 = (0)_2 \rightarrow 5 = (101)_2$$

$$17 = (10001)_2 \rightarrow 21 = (10101)_2$$

$$29 = (11101)_2 \rightarrow 29 = (11101)_2$$

# Encoding


Encode a standard deck of playing cards  
52 cards in 4 suits

How do we encode suits, face cards?

What operations do we want to make easy to implement?

Which is the higher value card?

Are they the same suit?



# Boolean Algebra

Developed by George Boole in 19th Century

Algebraic representation of logic (True  $\rightarrow$  1, False  $\rightarrow$  0)

AND:  $A \& B = 1$  when both A is 1 and B is 1

OR:  $A | B = 1$  when either A is 1 or B is 1

XOR:  $A \wedge B = 1$  when either A is 1 or B is 1, but not both

NOT:  $\sim A = 1$  when A is 0 and vice-versa

DeMorgan's Law:  $\sim (A | B) = \sim A \& \sim B$

$\sim (A \& B) = \sim A | \sim B$

AND		
$\&$	0	1
0	0	0
1	0	1

OR		
	0	1
0	0	1
1	1	1

XOR		
$\wedge$	0	1
0	0	1
1	1	0

NOT	
$\sim$	
0	1
1	0

# General Boolean Algebras

Operate on bit vectors

Operations applied bitwise

All of the properties of Boolean algebra apply

01101001	01101001	01101001	
<u>&amp; 01010101</u>	<u>  01010101</u>	<u>^ 01010101</u>	<u>~ 01010101</u>

Examples of useful operations:

$$x \wedge x = 0$$

	01010101
<u>^</u>	01010101
	00000000

$$x | 1 = 1, \quad x | 0 = x$$

	01010101
<u> </u>	11110000
	11110101

## Two possible representations

- 1) 1 bit per card (52): bit corresponding to card set to 1



low-order 52 bits of 64-bit word

“One-hot” encoding (similar to set notation)

Drawbacks:

Hard to compare values and suits

Large number of bits required

- 2) 1 bit per suit (4), 1 bit per number (13): 2 bits set



4 suits

13 numbers

Pair of one-hot encoded values

Easier to compare suits and values, but still lots of bits used

## Two better representations

- 3) Binary encoding of all 52 cards – only 6 bits needed  
 $2^6 = 64 \geq 52$



low-order 6 bits of a byte

Fits in one byte (smaller than one-hot encodings)

How can we make value and suit comparisons easier?

- 4) Separate binary encodings of suit (2 bits) and value (4 bits)







suit

value

Also fits in one byte, and easy to do comparisons

K	Q	J	...	3	2	A
1101	1100	1011	...	0011	0010	0001

	00
	01
	10
	11

# Compare Card Suits

```
char hand[5];           // represents a 5-card hand
char card1, card2;      // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if ( isSameSuit(card1, card2) ) { ... }
```

```
#define SUIT_MASK  0x30
```

```
int isSameSuit(char card1, char card2) {
    return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```

SUIT\_MASK = 0x30 =



# Compare Card Values

```
char hand[5];           // represents a 5-card hand
char card1, card2;      // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if ( greaterValue(card1, card2) ) { ... }
```

```
#define VALUE_MASK  0x0F

int greaterValue(char card1, char card2) {
    return ((unsigned char) (card1 & VALUE_MASK) >
            (unsigned char) (card2 & VALUE_MASK));
}
```

VALUE\_MASK = 0x0F =





# Encoding Integers

The hardware (and C) supports two flavors of integers

*unsigned* – only the non-negatives

*signed* – both negatives and non-negatives

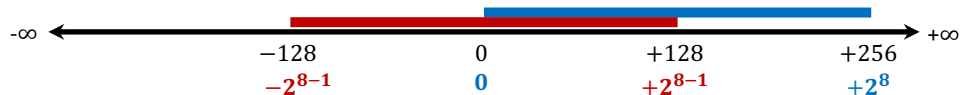
Cannot represent all integers with  $w$  bits

Only  $2^w$  distinct bit patterns

Unsigned values:  $0 \dots 2^w - 1$

Signed values:  $-2^{(w-1)} \dots 0 \dots 2^{(w-1)} - 1$

**Example:** 8-bit integers (e.g. `char`)



# Unsigned Integers

Unsigned values follow the standard base 2 system

$$b_7b_6b_5b_4b_3b_2b_1b_0 \\ = b_72^7 + b_62^6 + \dots + b_12^1 + b_02^0$$

Add and subtract using the normal “carry” and “borrow” rules, just in binary

63	00111111
+ 8	+00001000
<hr/> 71	<hr/> 01000111

Useful formula: N ones in a row =  $2^N - 1$

How would you make *signed* integers?

# Sign and Magnitude

Designate the high-order bit (MSB) as the “sign bit”

$\text{sign}=0$ : positive numbers;  $\text{sign}=1$ : negative numbers

Benefits:

Using MSB as sign bit matches positive numbers with unsigned

All zeros encoding is still = 0

Examples (8 bits):

$0x00 = 00000000_2$  is non-negative, because the sign bit is 0

$0x7F = 01111111_2$  is non-negative ( $+127_{10}$ )

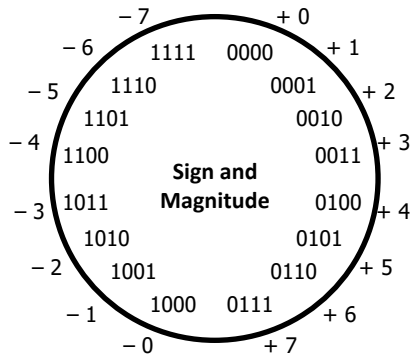
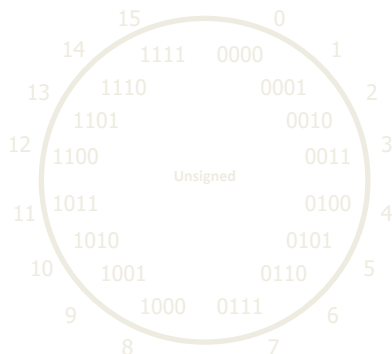
$0x85 = 10000101_2$  is negative ( $-5_{10}$ )

$0x80 = 10000000_2$  is negative...

# Sign and Magnitude

MSB is the sign bit, rest of the bits are magnitude

Drawbacks?

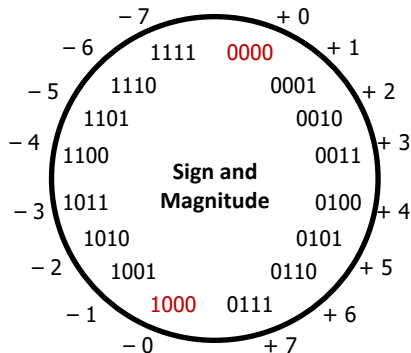


# Sign and Magnitude

MSB is the sign bit, rest of the bits are magnitude

Drawbacks:

Two representations of 0 (bad for checking equality)



# Sign and Magnitude

MSB is the sign bit, rest of the bits are magnitude

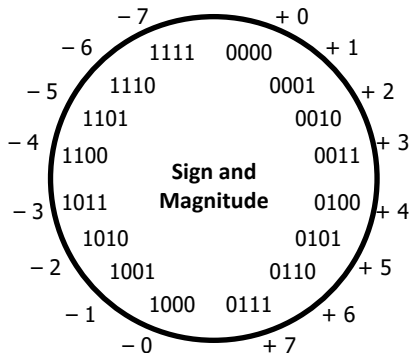
Drawbacks:

Two representations of 0 (bad for checking equality)

Arithmetic is cumbersome

Example:  $4 - 3 \neq 4 + (-3)$

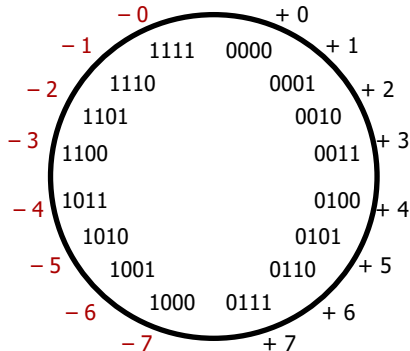
4	0100
<u>- 3</u>	<u>- 0011</u>
1	0001



# Two's Complement

Let's fix these problems:

- 1) "Flip" negative encodings so incrementing works



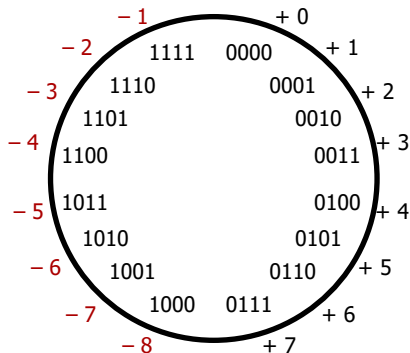
# Two's Complement

Let's fix these problems:

- 1) "Flip" negative encodings so incrementing works
- 2) "Shift" negative numbers to eliminate -0

MSB *still* indicates sign!

This is why we represent  
one  
more negative than positive  
number ( $-2^{N-1}$  to  $2^{N-1} - 1$ )





# Why Two's Complement is So Great

Roughly same number of (+) and (-) numbers

Positive number encodings match unsigned

Simple arithmetic ( $x + -y = x - y$ )


Single zero

All zeros encoding = 0

Simple negation procedure:

Get negative representation  
of any integer by taking  
bitwise complement and  
then adding one!

$( \sim x + 1 == -x )$



# Two's Complement Arithmetic

The same addition procedure works for both unsigned and two's complement integers

**Simplifies hardware:** only one algorithm for addition

**Algorithm:** simple addition, **discard the highest carry bit**

Called modular addition: result is sum *modulo*  $2^w$

## 4-bit Examples:

4    0100 +3    +0011 =7	-4    1100 +3    +0011 =-1	4    0100 -3    +1101 =1
--------------------------------	----------------------------------	--------------------------------

# Why Does Two's Complement Work?

For all integers  $x$ , we want:

$$\begin{array}{r} \textit{bit representation of } x \\ + \textit{ bit representation of } -x \\ \hline 0 \end{array} \quad (\text{ignoring the carry-out bit})$$


Start with an observation: what's  $x + \sim x$ ?

00000001	00000010	11000011
+ 11111110	+ 11111101	+ 00111100
<hr/>	<hr/>	<hr/>
11111111	11111111	11111111

So what's  $x + \sim x + 1$

# Signed/Unsigned Conversion Visualized

Two's Complement  $\rightarrow$  Unsigned  
Ordering Inversion  
Negative  $\rightarrow$  Big Positive



# Values To Remember

## Unsigned Values

$$\begin{aligned}\text{UMin} &= 0b00\dots0 \\ &= 0\end{aligned}$$

$$\begin{aligned}\text{UMax} &= \\ &0b11\dots1 \\ &= 2^w - 1\end{aligned}$$

## Two's Complement Values

$$\begin{aligned}\text{TMin} &= 0b10\dots0 \\ &= -2^{(w-1)}\end{aligned}$$

$$\begin{aligned}\text{TMax} &= 0b01\dots1 \\ &= 2^{(w-1)} - 1\end{aligned}$$

$$-1 = 0b11\dots1$$

## Example: Values for $w = 64$

	Decimal	Hex
UMax	18,446,744,073,709,551,615	FF FF FF FF FF FF FF FF
TMax	9,223,372,036,854,775,807	7F FF FF FF FF FF FF FF
TMin	-9,223,372,036,854,775,808	80 00 00 00 00 00 00 00
-1	-1	FF FF FF FF FF FF FF FF
0	0	00 00 00 00 00 00 00 00

# Arithmetic Overflow

Bits	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

When a calculation produces a result that can't be represented in the current encoding scheme

Integer range limited by fixed width  
Can occur in both the positive and negative directions

C and Java ignore overflow exceptions

You end up with a bad value in your program and no warning/indication...  
oops!

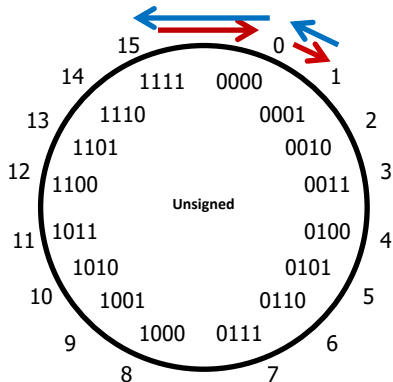
# Overflow: Unsigned

**Addition:** drop carry bit  
( $-2^N$ )

15	1111
+ 2	+ 0010
<hr/>	<hr/>
<del>17</del>	<del>10001</del>

**Subtraction:** borrow  
( $+2^N$ )

1	10001
- 2	- 0010
<hr/>	<hr/>
<del>-1</del>	1111
15	



# Overflow: Two's Complement

**Addition:**  $(+) + (+) = (-)$   
result?

$$\begin{array}{r} 6 \\ + 3 \\ \hline 9 \end{array}$$

~~9~~

$$\begin{array}{r} 0110 \\ + 0011 \\ \hline 1001 \end{array}$$

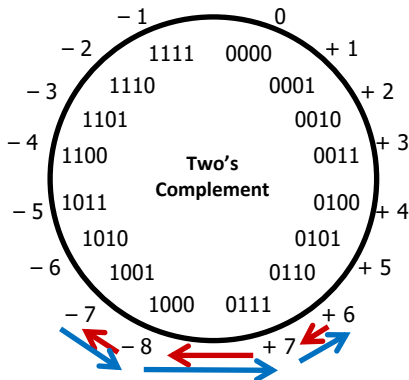
**Subtraction:**  $(-) + (-) = (+)$   
result?

$$\begin{array}{r} -7 \\ - 3 \\ \hline -10 \end{array}$$

~~-10~~

6

$$\begin{array}{r} 1001 \\ - 0011 \\ \hline 0110 \end{array}$$





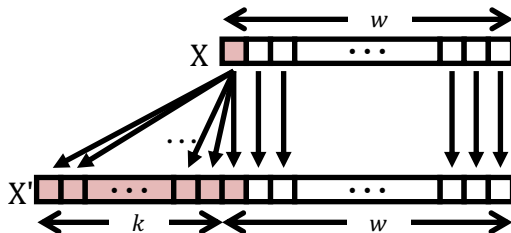
# Sign Extension

**Task:** Given a  $w$ -bit signed integer  $X$ , convert it to  $w+k$ -bit signed integer  $X'$  *with the same value*

**Rule:** Add  $k$  copies of sign bit

Let  $x_i$  be the  $i$ -th digit of  $X$  in binary

$$X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, \underbrace{x_{w-1}, x_{w-2}, \dots, x_1, x_0}_{\text{original } X}$$



# Sign Extension Example

Convert from smaller to larger integral data types  
C (and Java) automatically performs sign extension when converting to larger types

```
short int x =  
12345;  
int      ix = (int)  
x;  
short int y = -  
12345;  
int      iy = (int)  
y;
```

Var	Decimal	Hex	Binary
x	12345	30 39	00110000 00111001
ix	12345	00 00 30 39	00000000 00000000 00110000 00111001
y	-12345	CF C7	11001111 11000111
iy	-12345	FF FF CF C7	11111111 11111111 11001111 11000111

# Shift Operations

Left shift ( $x \ll n$ ) bit vector  $x$  by  $n$  positions

- Throw away (drop) extra bits on left

- Fill with 0s on right

Right shift ( $x \gg n$ ) bit-vector  $x$  by  $n$  positions

- Throw away (drop) extra bits on right

- Logical shift (for unsigned values)

  - Fill with 0s on left

- Arithmetic shift (for signed values)

  - Replicate most significant bit on left

  - Maintains sign of  $x$

# Shift Operations

Left shift ( $x \ll n$ )

Fill with 0s on right

Right shift ( $x \gg n$ )

Logical shift (for **unsigned** values)

Fill with 0s on left

Arithmetic shift (for **signed** values)

Replicate most significant bit on left

Notes:

Shifts by  $n < 0$  or  $n \geq w$  (bit width of  $x$ ) are *undefined*

**In C:** behavior of  $\gg$  is determined by compiler

depends on data type of  $x$  (signed/unsigned)

	$x$	0010 0010
	$x \ll 3$	0001 0 <b>000</b>
logical:	$x \gg 2$	<b>00</b> 00 1000
arithmetic:	$x \gg 2$	<b>00</b> 00 1000

	$x$	1010 0010
	$x \ll 3$	0001 0 <b>000</b>
logical:	$x \gg 2$	<b>00</b> 10 1000
arithmetic:	$x \gg 2$	<b>11</b> 10 1000

# Shifting Arithmetic?

What are the following computing?

$x \gg n$

0b 0100  $\gg 1$  = 0b 0010

0b 0100  $\gg 2$  = 0b 0001

Divide by  $2^n$

$x \ll n$

0b 0001  $\ll 1$  = 0b 0010

0b 0001  $\ll 2$  = 0b 0100

Multiply by  $2^n$

Shifting is faster than general multiply and divide operations

## Left Shifting Arithmetic 8-bit Example

No difference in left shift operation for unsigned and signed numbers (just manipulates bits)

Difference comes during interpretation:  $x * 2^n$ ?

			Signed	Unsigned
<code>x = 25;</code>	00011001	=	25	25
<code>L1=x&lt;&lt;2;</code>	0001100100	=	100	100
<code>L2=x&lt;&lt;3;</code>	00011001000	=	-56	200
<code>L3=x&lt;&lt;4;</code>	000110010000	=	-112	144


## Right Shifting Arithmetic 8-bit Examples

**Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values

Logical Shift:  $x / 2^n$ ?

`xu = 240u;`    11110000    = 240

`R1u=xu>>3;`    00011110000    = 30



`R2u=xu>>5;`    0000011110000    = 7


## Right Shifting Arithmetic 8-bit Examples

**Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values

**Arithmetic** Shift:  $x / 2^n$ ?

`xs = -16;`    11110000    = -16

`R1s=xs>>3;`    11111110    = -2



`R2s=xs>>5;`    11111111    = -1